

Data Mining project

Recommendation system for Medical Environment

Bertelli Davide
Artificial Intelligence Systems, 2nd year.
University of Trento
Trento, Italy
bertellidavide98@gmail.com

ABSTRACT

This report provides the algorithms, formulations and code implementations that allowed to build a recommendation system applicable to the medical environment. The algorithm leveraged on the biological similarity among the known patients to provide its suggestions. The system implemented is an hybrid one which uses a content-based rating method as a baseline for its ratings and improves them through a collaborative filtering algorithm. The results of the system had been evaluated using the RMSE scoring strategy over datasets containing different numbers of available patients, and leading to an average RMSE value of 0.31.

1 INTRODUCTION AND MOTIVATION

In a medical environment a patient asks for suggestions to a doctor regarding one or more diseases it is suffering from. Usually the therapies suggested came from the doctor own experience. In this perspective, the personal experience covers the highest importance in formulating a diagnosis of the disease and healing the patient. Given this environment it was thought to provide suggestions for a specific condition using as experience a dataset filled with the medical history of several patients. A good recommendation system for this scenario could serve as an artificial colleague for the medical staff or as a tool to allow the patients' self-treatment for the most common diseases, even though a professional opinion must be always required. What has been tried to achieve with this work is a Recommendation System able to suggest the best therapies to which a patient should undergo to cure a specific disease it is suffering from. To perform those suggestions the system must know the medical history of some patients, to use as experience, and the one of the patient under analysis, together with a list of therapies that may be suggested.

2 RELATED WORK

2.1 Recommendation Systems

The recommendation systems field is very extended, anyway it is possible to classify the available algorithms in three main categories:

- (1) The **Content-based Recommendation Systems** which work by providing suggestions based on the user's previous ratings for objects similar to the one to suggest. It has a dual implementation depending on which element it is more focused on. If more attention is given to the items available then the recommendation system will be an item-profile one,

otherwise if users are more relevant for the final application it will be a user-profile one.

- (2) The **Collaborative Filtering Recommendation Systems** which works by leveraging on the ratings of an object coming from several users. The idea behind this algorithm is that two users that have a similar taste may like a specific object in the same way. This system splits into two kinds: the model based and the memory based. The first exploits machine learning techniques for predicting the ratings, the latter uses the "neighbours" of the user-item combination under analysis to provide a prediction and it is further divided into item and user based collaborative filtering depending on which element to focus on.
- (3) The **Hybrid Recommendation Systems** are systems designed to provide more robust predictions merging methods of the previous categories or, even, combining their results.

Independently by its category, a Recommendation System, reaches a suggestion identifying which objects, users or items, are similar to the one under analysis or to the ones for which the ratings are available.

2.2 Similarity Scores

The similarity is a mathematical score that is attributed to pairs of objects. Its computation can be performed through different metrics such as:

- The **Jaccard Similarity** which is defined as the ratio of the intersection with the union. Meaning that two objects are similar if they have a lot of elements in common.

$$\frac{|obj_1 \cap obj_2|}{|obj_1 \cup obj_2|} \quad (1)$$

Where obj_i represents the ratings vector for a specific object.

- The **Euclidean Distance** which defines that two objects are similar if they are close to each other.

$$d = \sqrt{(obj_{1,1} - obj_{2,1})^2 + \dots + (obj_{1,N} - obj_{2,N})^2} \quad (2)$$

Where $obj_{i,j}$ represents the j^{th} rating of object i , Supposing that obj_1 and obj_2 are the ratings vectors associated to an object in an N dimensional space of users or items.

- The **Cosine similarity** which works by measuring the amplitude of the angle between the ratings vectors associated to two objects. The lesser the amplitude, the higher the similarity.

$$cos_{sim} = cos(obj_1, obj_2) = \frac{obj_1 * obj_2}{||obj_1|| ||obj_2||} \quad (3)$$

Where obj_i are the ratings vectors for some objects in the N dimensional space of user or items.

It exist a slight variation called **Centered Cosine Similarity** which is the same as defined above, but the vectors are previously normalized by subtracting their own mean.

2.3 Evaluation Metrics

Lastly, the Recommendation Systems need a tool able to measure their performances and to allow an objective comparison among different implementations. Some of the most used evaluation metrics are:

- The **K-fold** which works by splitting the data available into random subsets, folds, which will be used for training and testing: specifically k-1 for training and the left one for testing. The testing is performed many times with different folds and at the end the mean accuracy is computed as the average score of the system performances.
- The **Mean Squared Error MSE** is a metric computed by averaging the error occurred between the predicted ratings and their true values. It exist a variation called **Root Mean Squared Error** which is the square root of the MSE.

$$MSE = \frac{\sum_{t \in Test} |r_t - r_t^*|}{|Test|} \quad RMSE = \sqrt{MSE} \quad (4)$$

With r_t and r_t^* being the truth and predicted values, while $|Test|$ is the number of elements under testing.

2.4 Project Setup

In this work it has been built an Hybrid Recommendation System using the global baseline estimate method of the Content Based Recommendation Systems, to improve the prediction performances of a Collaborative Filtering one. The system built uses the euclidean distance to compute the similarity among the objects and it has been evaluated using the RMSE metric.

3 PROBLEM STATEMENT

Supposing an input dataset D composed of the recorded medical histories of some patients D_P , a set of known condition, D_C , and a set of available therapies D_T

$$D = D_P \cup D_T \cup D_C \quad (5)$$

the problem of suggesting the best therapies for a condition c to a patient P_q can be formally expressed in finding a utility function u such as:

$$u : (P_q|_c \cup D_P|_c) \times D_T \rightarrow r_t|_c \quad t \in \{D_T \setminus \{P_q^T|_c\}\} \quad (6)$$

Where: $P_q|_c$ is the set of rating scores for therapies that the patient P_q , the query patient, has already tried for healing itself from the condition c ; $D_P|_c$ is the set of rating scores for the therapies that the patients in the dataset have tried for the condition c ; D_T is the set of therapies available in the dataset; $P_q^T|_c$ is the set of therapies that have already been tried by the query patient; $r_t|_c$ is the rating of a therapy t for a condition c .

4 SOLUTION

The solution provided in this report, works as follows. Given in input the query patient P_q and the condition c for which

the suggestions have to be made, a utility matrix u_c is built parsing the medical history of all the patients in the given dataset, D_P , and retaining the rating scores of the therapies for those patients that underwent the specific condition.

$$u_c : \{P_i|_c\} \quad \forall P_i : P_i \in \{D_P|_c \cup P_q|_c\} \quad (7)$$

$$u_c = \begin{bmatrix} r_{1,1} & r_{1,2} & \dots & r_{1,j} \\ r_{2,1} & r_{2,2} & \dots & r_{2,j} \\ \vdots & \vdots & \vdots & \vdots \\ r_{i,1} & r_{i,2} & \dots & r_{i,j} \end{bmatrix}_{\substack{i \in \{D_P|_c\} \\ j \in \{D_T\}}} \quad (8)$$

Where $P_i|_c$ is the ratings vector of therapies that the patients in the dataset underwent in order to get cured by the condition c ; $D_P|_c$ is a subset of the input dataset containing the ratings of therapies for those patients which have been subjected to the condition c ; $P_q|_c$ is the ratings vector for the therapies already tried by the query patient for the condition c ; $r_{i,j}$ is the rating provided by the i^{th} patient to the j^{th} therapy, i.e. the rate of success.

Once the matrix is built, each rating vector in it undergo a weighting process dependent by its average goodness across the patients that used it:

$$r_{i,j} = \left(\frac{\sum_{i \in D_P|_c} r_{i,j}}{R_j * 100} \right) * r_{i,j} \quad \forall j \in D_T \quad (9)$$

Where R_j is the number of elements in the j^{th} therapy ratings vector. Then, the patient vectors are centered with respect to 0 by subtracting their mean value, in order to make all the scores comparable:

$$r_{i,j} = r_{i,j} - \left(\frac{\sum_{j \in D_T} r_{i,j}}{R_i} \right) \quad \forall i \in D_P|_c \quad (10)$$

Where R_i is the number of elements in the i^{th} patient ratings vector. Lastly, a normalization is performed in order to enforce the maximum value in the matrix to be 1.

$$r_{i,j} = \frac{r_{i,j}}{\max(u_t)} \quad \forall i \in D_P|_c \quad \forall j \in D_T \quad (11)$$

After the normalization steps, the query patient's ratings vector is removed from the utility matrix and it is used to find the most similar patients. As anticipated in section 2 the euclidean distance(2) is used to compute the similarity between patients and the most similar ones are identified by being the closest to the query patient. This metric is used instead of the cosine similarity due to the high sparsity of the utility matrix which negatively affected the computation of the cosine. To perform the computation of the distance, the empty cells of the matrix have been temporary filled with a 0 value. The nearest neighbours are identified as the patients within a distance less than half the length of the query patient vector. To further improve the similarity it has been thought to increase the similarity score for those patients which have a similar biology to the query one, as described in section 5.3. Then, for the patients found, their similarity score was risen in order to contribute more in the ratings prediction computation. Due to the usage of the euclidean distance, increase the similarity means to reduce the distance and so a -1 value was added to the similarity score of the patients having the same biology of the query one. Having found the nearest neighbours set, the prediction of the ratings for the therapies not tried

by the query patient begins. Specifically, it is possible to define the rating of a therapy t for the query patient P_q as:

$$r_{P_q,t} = b_{P_q,t} + \frac{\sum_{i \in NN} |1 - S_{P_q,i}| (r_{i,t} - b_{i,t})}{\sum_{i \in NN} |1 - S_{P_q,i}|} \quad (12)$$

Where $S_{p,i}$ is the similarity score between the query patient ratings vector and the ratings vector of the i^{th} patient in the Nearest Neighbours set NN . $r_{i,t}$ is the rating of the therapy t attributed by the i^{th} patient in the Nearest Neighbour set. $b_{p,t}$ and $b_{i,t}$ are the baseline estimates inherited by the Content based Recommendation systems. Their computation is performed using the following equations:

$$b_{p,t} = mean(u_c) + b_p + b_t \quad (13)$$

$$b_p = mean(P_q|_c) - mean(u_c) \quad (14)$$

$$b_t = mean(u_c^t) - mean(u_c) \quad (15)$$

Where $mean(u_c)$ is the mean value of the utility matrix, $mean(P_q|_c)$ is the mean value of the ratings vector of the query patient and $mean(u_c^t)$ is the mean value of the therapy given its ratings for each patient. It is important to notice the $|1 - S_{P_q,i}|$ in the computations which is due to the employment of the euclidean distance as the similarity metric. At the end of the computation the ratings vector of the query patient is completely filled with the estimates and the five highest scoring therapies, that have not already tried, are used as suggestions.

5 IMPLEMENTATION

5.1 Code Implementation

The Recommendation System has been written in python 3.8.10 and uses the pandas module[1] to manage the matrices and vectors. Custom data structures have been created to represent and better manage the data coming from the input dataset. The structures in use are:

- The **Entity** class which is used as the main structure which will be inherited by some of the others classes. It contains an identifier, a name and an optional type field.
- The **Condition** class which inherits the Entity one and is used to store the data necessary to represent a Condition in the dataset.
- The **Therapy** class which inherits the Entity one and is used to represent the data identifying a therapy in the dataset.
- The **Trial** class is the one used to represent a trial that a patient has done in order to heal itself from a disease. It contains a unique identifier, the begin and end dates of the trial, the condition identifier for which the trial was done and the rate of success.
- The **PCondition** class is the one used to store the data regarding the conditions affecting the patients. It contains its identifier, the diagnosed and cured dates and the identifier of the condition in the dataset.
- The **Patient** class which inherits the entity class and is used to manage the data representing the patients' medical history. Specifically, each patient is represented by a name in the format "name surname", an identifier and two lists: one containing the diseases that the patient have been subjected to and the other containing the trials done for each of them. This class

is especially used to represent the medical history of each patient.

- The **Dataset** class is the one used to represent the entire input dataset exploiting the classes defined so far. It is comprehensive of three lists representing the conditions or diseases, known from the input data, the therapies available in the dataset and the medical histories of the available patients. This class contains also the method which generates the utility matrix that will be further used to recommend the therapies.

5.2 Utility Matrix Construction

The utility matrix that is generated by the Dataset class is built creating a pandas.DataFrame object which rows and columns are respectively the patients and therapies identifiers gathered from the available objects in the Dataset class. Each cell of the matrix is filled by each rating that a patient gave for the therapies it tried. If a patient did not underwent a therapy and so there is not any rating available the NaN value is used as score. After the building task, a preprocessing step is performed in order to remove the columns, i.e. the therapies, without scores due to the inability to produce any rating for a therapy without anyone having used it.

5.3 Nearest Neighbours Computation

Then the utility matrix, exploiting the supported operations among the pandas objects, is subjected to the normalization operations described by the equations (9), (10) and (11). After the normalization steps the nearest neighbours of the query patients are computed using the euclidean distance (2) and saved into a DataFrame. Then each patient scorings vector in the utility matrix is filtered over the therapies that the query patient underwent and the scores are compared. If a match, with a 10% tolerance, is found, the patient is said to have a similar biology to the one of the query patient and its index is saved in a set. At this point the set of nearest neighbours patients is intersected with the one containing the indexes of patients having a similar biology. If the intersection is not empty the patients resulting from it have their euclidean distance decreased by 1 increasing the similarity and so giving more weight to their ratings during the prediction step.

5.4 Ratings Prediction

As already stated in section 4 the ratings prediction follows the equations (12), (13), (14) and (15) by exploiting the operations defined for the pandas objects. Once the ratings vector for the query patient is completely filled with the predictions, the therapies that it already tried for the condition under analysis are removed and the 5 highest scoring therapies are returned to the user inside a DataFrame object reporting the identifier, name, type and score of each selected therapy.

6 DATASET

The datasets used in the computations is randomly generated. It contains a fixed set of diseases and therapies plus a variable number of patients' medical histories.

6.1 Conditions and Therapies creation

The conditions and therapies available are taken directly from the internet by parsing two web pages[2][3] from which the names of the conditions and therapies are taken, while their types are inferred directly from the name. Once the couples name-type are generated, a new instance of the Condition or Therapy class is created and saved into two lists working as pools.

6.2 Patient Creation

The patient creation followed a completely different process: firstly, the names and surnames used for the *name* field have been sampled through the code made available by a github repository[4] which returned the most common names and surnames across the world. These were saved into two lists working as pools. Having at disposal also the pools for the therapies and conditions, a user-defined number of patients are created. After randomly sampling the names and surnames a unique identifier is generated. Then a randomly chosen number, between 0 and 10, is used to define the number of diseases in the medical history of the patient and for each of them the diagnosed and cured date are randomly generated. Each date is chosen by randomly sampling inside a precise interval, specifically the diagnosed date is sampled among the birth and death dates of the patient: the first is randomly chosen in an interval of 100 years going from the local date up to the past, the latter is generated by sampling a date in the interval going from the diagnosed date up to the death date which is a randomly chosen date computed by subtracting the randomly chosen age of the patient from 100(the expected human lifespan). The cured date is the only one allowed to be null, meaning that the disease is still not healed. Each condition has a probability of 20% to not being cured. After that, for each condition that the patient has, a random number of trials, between 0 and 10, is generated by randomly sampling the therapies to use from the respective pool. The dates of start and end of the treatments are chosen in the same way of the diagnosed and cured one, the only precaution taken is that for the same condition the trials have to be executed in a subsequent way.

7 EXPERIMENTAL EVALUATION

The recommendation system has been evaluated in terms of the Root Mean Squared Error(4) and the time it needs to generate a recommendation. The system has been built in a way that allows the user to provide the datasets used for evaluation or even resort to an automatic evaluation letting the code to generate the datasets. In this latter case the evaluation is run over four datasets containing different batches of patients: 25, 50, 100 and 150 thousands, while the number of conditions and therapies available are kept equal in all of them. The building of the random datasets follows the steps presented in section 6. Once the recommendation system has a list of datasets on which to work on, all of them are evaluated sequentially for a certain number of iterations defined by the user. In my experiments 30 iterations were used. The evaluation function firstly loads the datasets and saves their values inside the Dataset class, then it randomly choses a condition from the one saved in the class and uses it to build the utility

matrix. At this point it randomly samples from the indexes of the matrix a patient which has tried at least one therapy to cure the condition. Once the patient is chosen, a therapy is randomly sampled from the ones the patient underwent and the score of success is saved as the truth value, and it is substituted with an *NaN* value to allow the rating function to work properly. At this point, the nearest neighbours are computed and used together with the utility matrix, without the selected therapy value, to compute the ratings for the chosen therapy. Lastly the predicted value of the therapy and its real value are subjected to the RMSE computation(4) ending the first part of the evaluation process. The RMSE scores across the different iterations are collected and their mean value is used as the final RMSE score. During those steps the time requested by the algorithm to build the utility matrix, compute the nearest neighbours and predict the ratings for the therapy is collected and their mean is used as the final time score. The first evaluation of the system has been performed by running the evaluation algorithm over two distinct datasets of 100 thousands patients. The first dataset is an external one provided by the professor of the course, while the second is a randomly filled one. The results of the evaluation are reported in figures 1 and 2. As it




Figure 1: RMSE and time for the external dataset.




Figure 2: RMSE and time for the randomly filled dataset.

is possible to notice, for both datasets the recommendation system presented the same jittery behaviour with a mean RMSE of 0.4467 for the external dataset and 0.2028 for the randomly sampled one. The difference in the average may be due to the random sample at test time. Moreover the RMSE in the external dataset shows an higher mean value due to the lower dimensionality of therapies available, while in the randomly sampled one which has 204 dimensions the error influence may have a lower impact.

Regarding the execution time it is possible to observe that in the external dataset the computation time is halved with respect to the randomly generated one. This may be due to the number of therapies available in the datasets, specifically 51 in the external and 204 in the randomly filled one. The huge difference in the number of therapies may have influenced negatively the time of computation due to the vectors of ratings being 4 times bigger in the second dataset increasing the computation time. A second evaluation have been performed over the recommendation system aiming to analyze its performances in function of the number of patients available. In particular the analysis exploited four randomly filled dataset with 25, 50, 100 and 150 thousands patients.

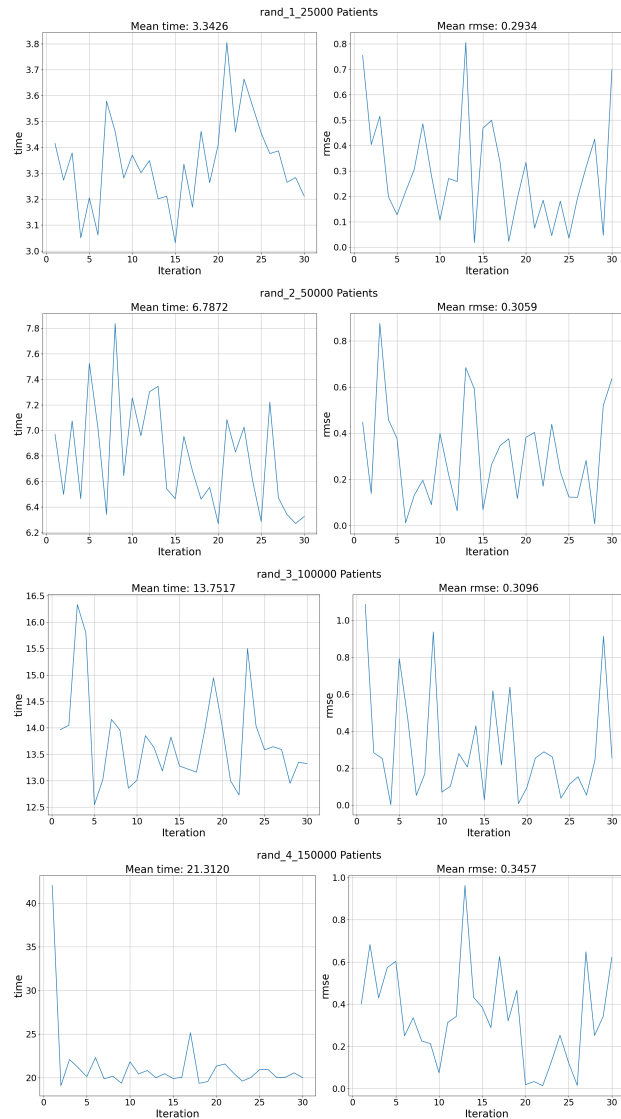


Figure 3: Results of the recommendation system over four datasets having 25, 50, 100 and 150 thousands patients




Figure 4: Average RMSE and time scores for the four randomly filled datasets.

As it is possible to notice from figure 4, the RMSE value is subjected to the stochasticity of the evaluation process and its values fluctuates, anyway it is possible to assert that its mean value is, as the previous evaluations, between 0.3136 with the highest peak at 0.35 and the lower one at 0.25. Generally speaking the best and worst cases may be seen as the extreme ones in which the worst is when the patient had done nothing to cure its disease and so the recommendation system resort to a content-base estimation, while the best one is when the patient had done some trials to get rid of its sickness and the recommendation system resort almost completely to the collaborative filtering component. It is possible to notice a constant increase in the RMSE with the number of patients in the dataset, anyway this behaviour may also be justified by the random nature of the datasets. About the time required by the system it is possible to infer a linearly proportional dependency with the number of patients, possibly due to the bigger utility matrix that has to be created for each disease to cure.

REFERENCES

- [1] <https://pandas.pydata.org/>
- [2] https://en.wikipedia.org/wiki/List_of_therapies
- [3] <https://www.nhsinform.scot/illnesses-and-conditions/a-to-z>
- [4] <https://github.com/philipperemy/name-dataset>